

# Implementation of Dynamic Trees with In-Subtree Operations

Tomasz Radzik

Department of Computer Science, King's College London  
Technical Report, March 1998

## ABSTRACT

We describe an implementation of dynamic trees with “in-subtree” operations. Our implementation follows Sleator and Tarjan’s framework of dynamic-tree implementations based on splay trees. We consider the following two examples of “in-subtree” operations. (a) For a given node  $v$ , find a node with the minimum key in the subtree rooted at  $v$ . (b) For a given node  $v$ , find a random node with key  $X$  in the subtree rooted at  $v$  (value  $X$  is fixed throughout the whole computation). The first operation may provide support for edge deletions in the dynamic minimum spanning tree problem. The second one may be useful in local search methods for degree-constrained minimum spanning tree problems. We conducted experiments with our dynamic-tree implementation within these two contexts, and the results suggest that this implementation may lead to considerably faster codes than straightforward approaches do.

## 1. INTRODUCTION

A *dynamic tree collection* (or simply *dynamic trees*) is a data structure which maintains a collection of node-disjoint rooted trees under a sequence of operations including the following structural updates.

- `make_tree( $k$ ):` Create a new node with key  $k$ . The created node forms a new single-node tree.
- `link( $v, w$ ):` Add an edge from node  $v$  to node  $w$  to join two trees into one. Precondition: nodes  $v$  and  $w$  are in different trees and  $v$  is the root of its tree.
- `cut( $v$ ):` Delete the edge from node  $v$  to its parent. This operation splits one tree into two trees, if node  $v$  has not been the root of its tree.
- `evert( $v$ ):` Make node  $v$  the root of its tree.

The applications which use dynamic trees usually require the following queries about the current structure of the trees.

- `find_root( $v$ ):` Find the root of the tree containing node  $v$ .

---

Address: Department of Computer Science, King’s College London, Strand, London WC2R 2LS, UK. Tel: +44-(0)171-873-2588. Fax: +44-(0)171-873-2851. E-mail: [radzik@dcs.kcl.ac.uk](mailto:radzik@dcs.kcl.ac.uk)  
The preliminary version of this paper was titled “Implementations of dynamic tree collections based on splay trees” and appeared in the *Proceedings of the Workshop on Algorithm Engineering, WAE ’97*, September 11-13, 1997, <http://www.dsi.unive.it/~wae97/proceedings>.

**parent( $v$ ):** Find the parent of node  $v$ .

For each node  $v$ , the parent of  $v$ , the children of  $v$ , and the subtree rooted at  $v$  are always defined with respect to the current root of the tree containing node  $v$ .

Each node  $v$  has a key, denoted by  $\text{key}(v)$ , and different choices of operations on the node keys define different variants of dynamic trees. Sleator and Tarjan [1985] assume that node keys are real numbers and consider the following operations.

**find\_key( $v$ ):** Return the value of the key of node  $v$ .

**min\_on\_path( $v$ ):** Find a node with the minimum key on the path from node  $v$  to the root.

**add\_key( $v, x$ ):** Add number  $x$  to the key of every node on the path from node  $v$  to the root.

Operations like **min\_on\_path** and **add\_key** refer to the path from a given node to the root, so we call them “on-path” operations.

Sleator and Tarjan [1985] show an implementation of the dynamic trees with the above two “on-path” operations (see also Tarjan [1983]). Their data structure is based on self-adjusting binary trees called *splay trees*, and requires  $O(\log n)$  *amortized* time per one dynamic-tree operation. We denote by  $n$  the maximum number of nodes in one tree over a given sequence of dynamic-tree operations. The dynamic trees with operations **min\_on\_path** and **add\_key** find applications in various network optimisation problems. For example, it has been shown that Sleator and Tarjan’s implementation of such dynamic trees improves asymptotic bounds on the running times of many old as well as newly proposed network flow algorithms [Ahuja et al. 1989; Goldberg et al. 1990; Goldberg and Tarjan 1989; Goldberg and Tarjan 1990; Sleator and Tarjan 1983]. Their data structure can also be used in the dynamic minimum spanning tree problem to support insertions of new edges in  $O(\log n)$  amortized time per one insertion [Frederickson 1985]. It should be noted that network flow algorithms and edge insertions in dynamic minimum spanning tree algorithms actually require dynamic trees with keys associated with edges not with nodes. However, in the context of such algorithms, the problem of maintaining edge keys can be easily reduced, without affecting asymptotic bounds on the running times, to the problem of maintaining node keys. Sleator and Tarjan [1983] describe an implementation of the dynamic trees with “on-path” operations which is based on biased search trees [Bent et al. 1985] and requires  $O(\log n)$  *worst-case* time per one dynamic-tree operation, but has a practical disadvantage of being considerably more complicated than the data structure based on splay trees.

In this paper we consider dynamic trees in which the queries involving node keys refer to subtrees rather than to paths. We consider the following three types of such queries.

**find\_random\_X( $v$ ):**

Node keys are drawn from an arbitrary set. One key value, called value **X**, is distinguished. Operation **find\_random\_X( $v$ )** selects a random node with key **X** in the subtree rooted at node  $v$ .

`min_in_subtree1(v)`:

Node keys are real numbers. Operation `min_in_subtree1(v)` finds a node with the minimum key in the subtree rooted at node  $v$ .

`min_in_subtree2(v, A)`:

Each node has  $q$  real-number keys. Number  $q$  does not have to be a constant, but is fixed for the whole computation. The  $i$ -th key of a node  $v$  is denoted by  $\text{key}(v, i)$ . Operation `min_in_subtree2(v, A)`, where  $v$  is a node and  $A$  is an array of size  $q$ , finds a node  $x_{\min}$  and an index  $i_{\min}$  which minimise  $\text{key}(x, i)$  over all nodes  $x$  in the subtree rooted at node  $v$  and all indices  $i$  such that  $A[i] = 1$ . That is, array  $A$  contains the characteristic vector of the set of indices of keys which are to be considered.

The applications of dynamic trees with “in-subtree” queries considered in this paper require also operation `new_key`, which changes a given node key to a given value. We describe an implementation of dynamic trees which gives  $O(\log n)$  amortized time per one operation, if operation `find_random_X` or `min_in_subtree1` is supported, and  $O(q \log n)$  amortized time per one operation, if operation `min_in_subtree2` is supported. Our implementation follows Sleator and Tarjan’s implementation of dynamic trees with “on-path” operations based on splay trees. The added difficulty of implementing “in-subtree” operations as opposed to “on-path” operations is that in order to search efficiently a tree down towards the leaves, we need some mechanism of selecting the appropriate child from the possibly many children of the currently considered node. On the other hand, searching the path from a given node to the root seems to be a more straightforward process since each node has at most one parent. In Section 2 we describe details of Sleator and Tarjan’s framework for implementing dynamic trees using splay trees, and in Section 3 we show how one can incorporate “in-subtree” operations into this framework.

Dynamic trees which support operation `find_random_X` can be used in implementations of local-search methods for constrained minimum spanning tree problems. Consider, for example, the NP-hard problem of finding a minimum-weight degree-3 spanning tree in a given weighted complete graph. A degree-3 tree is a tree such that each node has degree at most 3. If we apply the simulated-annealing method [Cerny 1985; Kirkpatrick et al. 1983] to this problem, we may have to maintain a degree-3 spanning tree under a sequence of operations of the following type. Remove a random edge from the current degree-3 spanning tree and reconnect the two parts into a new degree-3 spanning tree by adding a random edge which links two nodes of degree at most 2. Let the current degree-3 spanning tree  $T$  be rooted at an arbitrary node, and let each node  $v$  have either key  $X$  or key `not_X`, depending whether the degree of  $v$  in  $T$  is less than 3 (the node is available) or equal to 3 (the node cannot accept any additional tree edge). The described operation of exchanging random edges can be implemented by a constant number of dynamic-tree operations including operation `find_random_X`. In Section 4 we discuss details of such an implementation and its actual performance.

Dynamic trees with operations `min_in_subtree1` and `min_in_subtree2` can be used in the dynamic minimum spanning tree problem to support deletions of tree edges (that is, edges which are in the current minimum spanning tree). Let  $G$  denote

the underlying dynamic undirected weighted graph, and let  $T$  denote the current minimum spanning tree of  $G$  (assume that  $G$  is always connected). For each node  $v$  in  $G$ , let  $T_v$  be a copy of tree  $T$ . The key of a node  $x$  in a tree  $T_v$  is equal to the weight of edge  $(v, x)$  in graph  $G$ , or to infinity, if  $(v, x)$  is not an edge in  $G$ . Each tree  $T_v$  is rooted at an arbitrary node. If a tree edge is deleted from the graph, tree  $T$  breaks into two trees  $T'$  and  $T''$ , and we have to find a replacement edge which reconnects these two trees back into a minimum spanning tree. Each node  $v$  contributes one candidate for the replacement edge: a minimum-weight edge which is adjacent to node  $v$  and joins trees  $T'$  and  $T''$ . Such an edge can be found by applying to tree  $T_v$  a constant number of dynamic-tree operations, including operation `min_in_subtree1`. Following this approach, we implement deletions of tree edges so that one deletion takes  $O(n \log n)$  amortized time. Our experiments suggest that for dense graphs such an implementation can be quite competitive in practise. As discussed in Section 5, methods with considerably better asymptotic bounds have already been known for some time. However, consistently fast codes based on those asymptotically fast methods are yet to be developed. The existing codes sometimes perform substantially worse than even the ad-hoc method of re-computing the minimum spanning tree from scratch whenever a tree edge is deleted (see the experimental results reported by Amato, Cattaneo, and Italiano [1997]). In Section 5 we describe in detail how one can apply operations `min_in_subtree1` and `min_in_subtree2` to the dynamic minimum spanning tree problem, and we also present results from experiments with our dynamic minimum spanning tree codes.

All our codes are written in C++. Code `mst_recompute` is based on LEDA [Mehlhorn and Näher 1995]. Edge insertions in all our dynamic minimum spanning tree codes are supported by the implementation of Sleator and Tarjan's dynamic trees (with edge keys) developed by Sorin Moise. All experiments were conducted on Sun UltraSPARC ( $2 \times 300$  MHz, 512 MB). Our codes are available at <http://www.dcs.kcl.ac.uk/staff/radzik>.

## 2. GENERIC DYNAMIC TREES FOR IN-SUBTREE OPERATIONS

Our implementation of dynamic trees with “in-subtree” operations is based on the representation of dynamic trees by *virtual trees* proposed by Sleator and Tarjan [1985].

### 2.1 Virtual trees

A *virtual tree*  $D$  is a rooted tree such that for each node  $v$  in  $D$ , each child of  $v$  has one label from the set  $\{\text{left}, \text{right}, \text{dash}\}$ . A node cannot have more than one left and one right child, but may have a number of dash children. The left and the right children of a node  $v$  are called the *solid children* of  $v$ . Edges between nodes and their solid children are called *solid edges*, while the other edges are called *dash edges*. A maximal binary tree in  $D$  formed by solid edges is called a *solid tree*, and its root is called a *solid root*. For a node  $v$  of a virtual tree, `solid_root(v)` denotes the root of the solid tree containing  $v$ .

Each node  $v$  of a virtual tree has a *reverse bit*  $b(v)$  associated with it. Let  $B(v)$  denote the parity of the sum of the reverse bits on the path from node  $v$  to `solid_root(v)`. A solid child  $u$  of node  $v$  is the *true left child* of  $v$ , if  $u$  is the left child of  $v$  and  $B(v) = 0$ , or  $u$  is the right child of  $v$  and  $B(v) = 1$ . Otherwise node  $u$  is

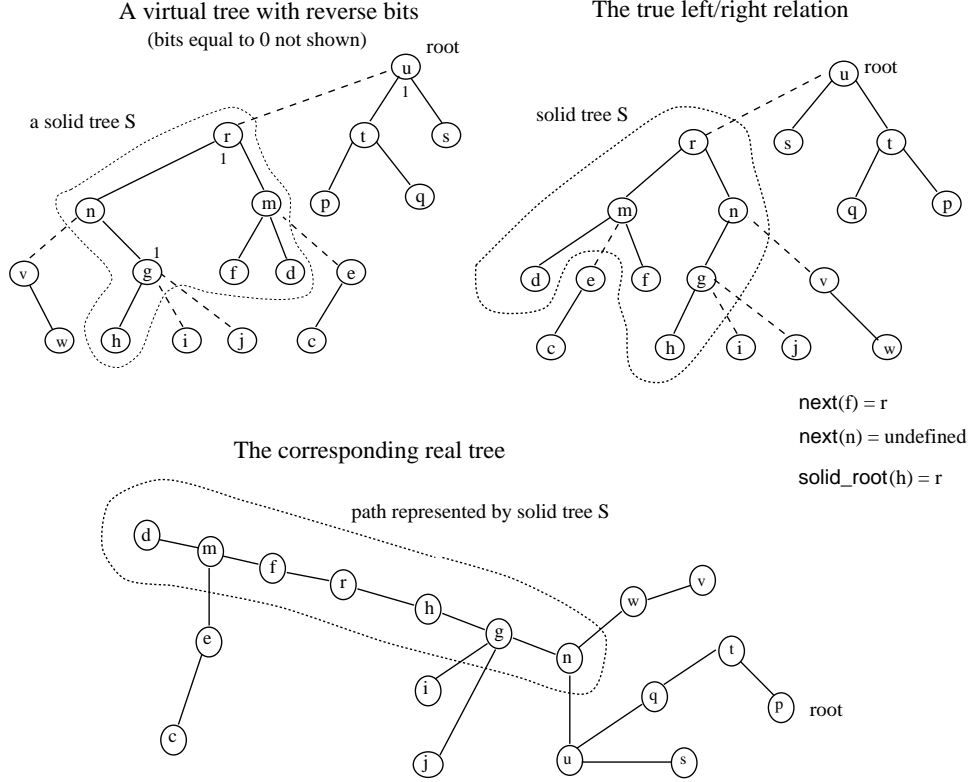


Fig. 1. A virtual tree, the true left-right relation in this tree, and the corresponding real tree.

the *true right child* of  $v$ . The *true infix order* of the nodes of a solid tree is the infix order with respect to the true left-right relation. For a node  $v$  of a virtual tree  $D$ , let  $\text{next}_D(v)$  denote the successor of node  $v$  in the true infix order of the nodes of the solid tree containing node  $v$ .

Throughout this paper we use adjective “real” when we want to make it clear that we refer not to virtual trees but to dynamic trees which are represented by virtual trees. A virtual tree  $D$  represents a real tree  $T$ , if and only if, both trees have the same set of nodes and for each node  $v$ ,

$$\text{parent}_T(v) = \begin{cases} \text{next}_D(v), & \text{if } \text{next}_D(v) \text{ exists,} \\ \text{parent}_D(\text{solid\_root}_D(v)), & \text{if } \text{next}_D(v) \text{ does not exist.} \end{cases}$$

The introduced terminology and the correspondence between a virtual tree  $D$  and the real tree  $T$  represented by  $D$  are illustrated in Figure 1. Observe that the solid trees in tree  $D$  correspond to some paths in tree  $T$ . Actually, the main concept of virtual trees is to provide a mechanism for representing long paths appearing in dynamic trees by binary trees with possibly small depths.




Fig. 2. Operation **rotate** in a virtual tree. “0” indicates that the reverse bit must be 0.

## 2.2 Primitive operations on virtual trees

We identify the following primitive operations on virtual trees: **unreverse**, **rotate**, **splice**, **connect**, **disconnect**, **switch\_bit**, and **make\_tree**. Operations **unreverse**, **rotate**, and **splice** have one argument, a node  $v$ , and alter locally the virtual tree  $D$  containing node  $v$  but do not affect the real tree represented by  $D$ . Operation **unreverse**( $v$ ) establishes at node  $v$  the locally correct left-right relation by performing the following update. If  $b(v) = 1$ , then  $b(v)$  is set to 0, the solid children of  $v$  are swapped, and their reverse bits are reversed. Operation **rotate**( $v$ ) rearranges locally the subtree rooted at the solid parent  $u$  of node  $v$  as shown in Figure 2 (this is the rotation operation used in binary search trees). Node  $v$  moves one level up the tree and becomes the parent of node  $u$ . The precondition of this operation is that  $b(v) = b(u) = 0$ . Operation **splice**( $v$ ) can be applied when node  $v$  is a dash child of a solid root  $u$  and  $b(u) = 0$ . This operation swaps node  $v$  with its left sibling (see Figure 3).

Primitive virtual-tree operations **connect**, **disconnect**, and **switch\_bit** affect the real trees. Operation **connect**( $v, label, u$ ), where  $label \in \{left, right, dash\}$ , makes node  $v$  a child of node  $u$  as indicated by  $label$ . The required precondition is that nodes  $v$  and  $u$  are in different virtual trees, node  $v$  is the root of its virtual tree, and node  $u$  does not have the left (right) child, if  $label = left$  ( $label = right$ ). Operation **disconnect**( $v$ ) cuts the connection between node  $v$  and its parent. Operation **switch\_bit**( $v$ ) changes the value of the reverse bit at node  $v$  to the opposite one (and, consequently, reverses the infix order of the nodes of the solid subtree rooted at  $v$ ).

The virtual-tree operation **make\_tree** creates a new single-node virtual tree. Since a single-node dynamic tree  $T$  is essentially the same as the single-node virtual tree  $D$  representing  $T$ , the dynamic-tree operation **make\_tree** is simply the virtual-tree operation **make\_tree** plus appropriate initialisation of the key-related attributes of the new node.

We do not assume that each primitive virtual-tree operation takes constant time. For example, operation **rotate** may involve some updates of key-related attributes, and these updates may require more than constant time.

## 2.3 Operation splay

Let  $v$  be a node in a virtual tree  $D$ , let  $u$  be the root of  $D$ , and let  $w$  be the root of the solid tree  $S$  containing node  $v$ . Operation **solid\_splay**( $v$ ) rearranges tree  $S$  so that node  $v$  becomes the root of  $S$ . The precondition of this operation is that all reverse bits on the path in  $D$  from  $v$  to  $w$  are equal to 0. The rearrangement of tree  $S$  is accomplished by some sequence of operations **rotate** applied to nodes on the path from  $v$  to  $w$ . Operation **solid\_splay**( $v$ ) is actually Sleator and Tarjan’s splaying




Fig. 3. Operation `splice` in a virtual tree. The reverse bit of node  $u$  must be 0.

operation in binary search trees [Sleator and Tarjan 1985] applied to node  $v$  of tree  $S$ .

Operation `splay`( $v$ ) brings node  $v$  to the root of the virtual tree by the following computation. First we apply operation `unreverse` to each node on the path from the root  $u$  to node  $v$ . Then we walk up the virtual tree from  $v$  to  $u$  performing operation `solid_splay` in each encountered solid tree. The path from node  $v$  to the (new) root  $z$  of the virtual tree consists now only of dash edges. Next we perform operation `splice` at each node on the path from  $v$  to  $z$ , and finally apply operation `solid_splay` to node  $v$ . Observe that operation `splay` uses only primitive virtual-tree operations `unreverse`, `rotate`, and `splice`, so it does not affect the real tree represented by virtual tree  $D$ . After operation `splay`( $v$ ) has been executed, the reverse bit of the new root (node  $v$ ) is equal to 0. We use this condition in designing implementations of dynamic-tree operations.

Sleator and Tarjan [1985] show that any sequence of  $m$  virtual-tree operations `splay`, `make_tree`, `connect`, `disconnect`, and `switch_bit`, which is applied to the empty collection of virtual trees and includes  $n$  operations `make_tree`, requires  $O(m \log n)$  primitive virtual-tree operations.

#### 2.4 Implementation of basic dynamic tree operations

Operation `splay` is the main tool in implementations of dynamic-tree operations. This operation plays a dual role: It makes the crucial nodes roots of virtual trees, so necessary updates of the attributes of these nodes become straightforward; and it implicitly balances virtual trees, so the amortized running time of one dynamic-tree operation is low. Implementations of the basic dynamic-tree operations `link`, `cut`, and `evert`, are shown below.

```

cut( $v$ ):    splay( $v$ );
           if  $v$  has the right child  $w$  then disconnect( $w$ )
           else  $v$  is the root of its real tree (so there is nothing to cut).

link( $v, w$ ): splay( $v$ ); splay( $w$ );
           if  $w$  does not have the left child then connect( $v, left, w$ )
           else connect( $v, dash, w$ ).

evert( $v$ ):  splay( $v$ );
           if  $v$  has the right child  $w$  then
               disconnect( $w$ ); switch_bit( $w$ );
               if  $v$  does not have the left child then connect( $w, left, v$ )
               else connect( $w, dash, v$ )
           else  $v$  is the root of its real tree.

```




Fig. 4. Operation  $\text{evert}(v)$ . (a) The virtual tree after operation  $\text{splay}(v)$  and the represented real tree (the arrows point towards the root). (b) The virtual tree and the represented real tree after the updates of operation  $\text{evert}(v)$ . Trees  $Ar$ ,  $Br$ , and  $Cr$  are virtual trees representing real trees  $A$ ,  $B$ , and  $C$ .

Let us, for example, examine the implementation of operation  $\text{evert}(v)$ . We assume that the edges of a real tree are directed towards the root. To make node  $v$  the root of its real tree, we have to reverse the directions of all edges on the path  $P$  from  $v$  to the root. Operation  $\text{evert}(v)$  begins with applying operation  $\text{splay}$  to node  $v$ . Node  $v$  becomes the root of its virtual tree and the solid right subtree  $S$  of node  $v$  represents path  $P$  (see Figure 4(a)). We reverse the directions of the edges on this path by switching the reverse bit at the root of subtree  $S$  and making  $S$  the left, if possible, or a dash subtree of  $v$  (see Figure 4(b)).

Query  $\text{find\_root}(v)$  is implemented in the following way. We first apply operation  $\text{splay}$  to node  $v$  to bring  $v$  to the root of the virtual tree. The root  $r$  of the real tree is now the last node in the solid right subtree of  $v$ , according to the true infix order of nodes. Once we locate node  $r$ , we finish the computation by executing operation  $\text{splay}(r)$ . Such “final” operations  $\text{splay}$  are required to maintain a desired bound on the amortized running time of one dynamic-tree operation.

```

find_root(v):  splay(v);  r ← v;
               while r has the right child do
                 r ← the right child of r;  unreverse(r);
               splay(r);  return r.

```

The implementation of query  $\text{parent}(v)$  is similar to the implementation of query  $\text{find\_root}(v)$ . The difference is that after operation  $\text{splay}(v)$ , we search not for the last but for the first node in the solid right subtree of  $v$ .

A sequence of  $m$  dynamic-tree operations, which is applied to the empty collection of dynamic trees and includes  $n$  operations  $\text{make\_tree}$ , is implemented by a



sequence of  $O(m)$  virtual-tree operations `splay`, `make_tree`, `connect`, `disconnect`, and `switch_bit`, which requires  $t_1 = O(m \log n)$  primitive virtual-tree operations, plus some additional interstitial computation, which takes  $t_2 = O(t_1)$  time. Thus the amortized time of one dynamic-tree operation is  $O(d \log n)$ , where  $d$  is a bound on the running time of one primitive virtual-tree operation.

### 3. IMPLEMENTATION OF “IN-SUBTREE” QUERIES

Let us consider the dynamic-tree operation `min_in_subtree1(v)`. Each node  $w$  of a dynamic tree has a real number `key(w)` associated with it. We want to find a node with the minimum key in the subtree rooted at node  $v$ . To be able to locate quickly such a node, we maintain at each node  $w$  of a virtual tree the key `key(w)` and the minimum key `min_key(w)` in the virtual subtree rooted at  $w$ .

Operation `min_in_subtree1(v)` begins with executing operation `splay(v)`. Node  $v$  is now the root of its virtual tree  $D$ , and the set of nodes of the real subtree rooted at  $v$  is the subset of nodes of the virtual tree  $D$  which consists of node  $v$ , the nodes of the left subtree of  $v$ , and the nodes of all dash subtrees of  $v$ . We navigate from the root  $v$  down the virtual tree  $D$  using values `min_key`. When we located a node  $x$  with the minimum key, we complete the computation by applying operation `splay` to node  $x$  (the “final” `splay`). Note that values `min_key` have to be appropriately updated during each primitive virtual-tree operation.

Let  $k$  denote the maximum degree of a node of a real tree. If the dash children of each node of a virtual tree are maintained simply in an unordered list, then one primitive virtual-tree operation may require  $O(k)$  time. This would imply that the amortized time of one operation `splay`, and consequently the amortized time of one dynamic-tree operation, is  $O(k \log n)$ . To achieve an  $O(\log n)$  bound, we use virtual trees to represent directly only such dynamic trees that each node has at most two children. We call such trees *restricted dynamic trees*. We implement *unrestricted dynamic trees* by representing them by restricted dynamic trees.

#### 3.1 Restricted dynamic trees

We represent restricted dynamic trees by virtual trees such that each node has at most one dash child. Thus the information stored at a virtual node  $v$  consists of:

- (a) four pointers, which point to the (virtual) parent, the left child, the right child, and the dash child of node  $v$ ,
- (b) the reverse bit of node  $v$ , and
- (c) the information regarding the node keys.

The information regarding the node keys depends on the type of keys and the type of operations we want to supported. To support operation `min_in_subtree1`, each node  $v$  of a virtual tree contains `key(v)` and `min_key(v)`, as described above. To support operation `min_in_subtree2`, each node  $v$  contains an array `key` of  $q$  keys of node  $v$ , and an array `min_key` of  $q$  minimum keys in the virtual subtree rooted at  $v$ : `key(v, i)` denotes the  $i$ -th key of node  $v$  and `min_key(v, i)` is equal to the minimum of `key(x, i)` over all nodes  $x$  in the virtual subtree rooted at  $v$ . To support operation `find_random_X`, we store at each node  $v$  the key `key(v)` of node  $v$  and the number of keys  $X$  in the virtual subtree rooted at  $v$ .

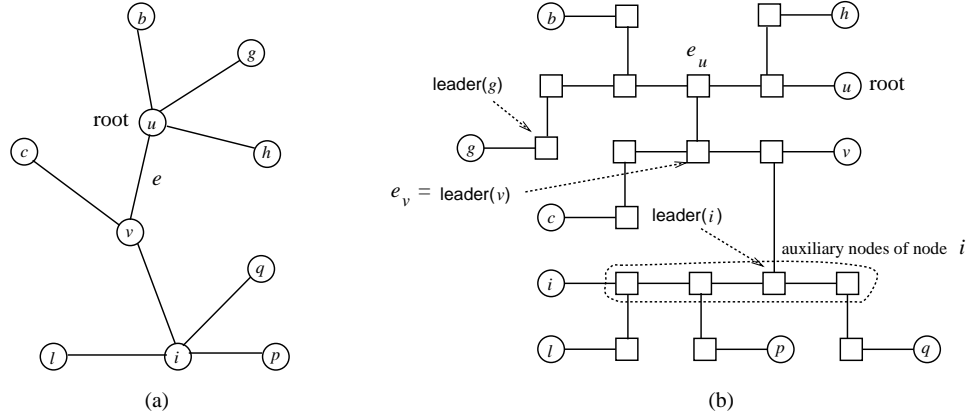


Fig. 5. (a) Unrestricted real tree  $T$ . (b) The corresponding restricted tree  $R$ .

In such virtual trees, each primitive virtual-tree operation runs in constant time, if we implement operation `min_in_subtree1` or `find_random_X`, and in  $O(q)$  time, if we implement operation `min_in_subtree2`. Thus in these two cases, the amortized time of one dynamic-tree operation on restricted trees is  $O(\log n)$  and  $O(q \log n)$ , respectively.

### 3.2 Unrestricted dynamic trees

We represent unrestricted dynamic trees by restricted ones in the following way. Let  $T$  be an unrestricted dynamic tree. The restricted tree  $R$  representing tree  $T$  is essentially the adjacency-list structure of tree  $T$  (see Figure 5). Tree  $R$  has two types of nodes: the real nodes, that is, the nodes of tree  $T$ , and auxiliary nodes, which are used to represent the edges of tree  $T$ . If  $e = (v, u)$  is an edge in tree  $T$ , then there are two auxiliary nodes  $e_v$  and  $e_u$  and edge  $(e_v, e_u)$  in tree  $R$ . The auxiliary node  $e_v$  is associated with the real node  $v$ . A real node  $v$  in tree  $R$  and all its auxiliary nodes are connected by edges to form a path; node  $v$  is an end node of this path. The root of tree  $R$  is the same as the root of tree  $T$ . If  $e = (v, u)$  is an edge in tree  $T$  and  $u$  is the parent of  $v$ , then the auxiliary node  $e_v$  in tree  $R$  is called the *leader* of node  $v$  in tree  $R$  (see Figure 5). Observe that the real nodes of the subtree of tree  $R$  rooted at the leader of a real node  $v$  are exactly the nodes of the subtree of tree  $T$  rooted at  $v$ .

Some dynamic-tree operations on a collection  $C_1$  of unrestricted trees, for example operations `find_root` and `evert`, are simply the same operations on the collection  $C_2$  of restricted trees representing collection  $C_1$  (that is, for example, operation `evert(v)` on  $C_1$  is implemented by operation `evert(v)` on  $C_2$ ). The other dynamic-tree operations on unrestricted trees can be implemented by a constant number of dynamic-tree operations on restricted trees. Actually, to achieve this we need two additional operations on restricted trees: `childr` and `leader`. We use subscripts “r” and “u,” when we want to make it clear whether we refer to restricted or unrestricted dynamic trees. For a node  $x$  of a restricted dynamic tree  $R$ , operation `childr(x)` returns an arbitrary child of node  $x$  in tree  $R$ . The implementation of this




Fig. 6. (a) Operation  $\text{cut}(v)$  in an unrestricted tree. (b) The update of the corresponding restricted tree.

operation is based on the observation that after executing operation  $\text{splay}(x)$ , the children of node  $x$  in tree  $R$  are the rightmost nodes in the left and the dash virtual subtrees of node  $x$ .

Operation  $\text{leader}(v)$  applies to a real node  $v$  of a restricted dynamic tree  $R$  representing an unrestricted dynamic tree, and returns the leader of node  $v$  in tree  $R$ . If node  $v$  is the root of tree  $R$ , then  $\text{leader}(v)$  returns  $v$ . To implement this operation, we maintain at each node  $x$  of tree  $R$  an attribute  $\text{real\_node}(x)$ . If  $w$  is a real node in tree  $R$  and  $z$  is an auxiliary node of node  $w$  in tree  $R$ , then  $\text{real\_node}(z) = \text{real\_node}(w) = w$ . Let  $P$  be the path in tree  $R$  from node  $v$  to the root. The implementation of operation  $\text{leader}(v)$  is based on the observation that the nodes on path  $P$  with attribute  $\text{real\_node}$  equal to  $v$  form an initial segment of  $P$ , and the leader of node  $v$  is the last node of this segment. Thus to find the leader of node  $v$ , perform operation  $\text{splay}(v)$  and then find in the solid tree rooted at node  $v$  the last node  $x$ , in the true infix order, such that  $\text{real\_node}(x) = v$ .

Operation  $\text{parent}$  on unrestricted dynamic trees has the following implementation.

$\text{parent}_u(v) : \text{return } \text{real\_node}(\text{parent}_r(\text{leader}(v)))$ .

To implement the “in-subtree” operations, we set the keys of auxiliary nodes in such a way that they do not interfere with the keys of the real nodes. If we implement the minimum-in-subtree queries, we set the keys of auxiliary nodes to infinity. If we implement operation  $\text{find\_random\_X}$ , we set the keys of auxiliary nodes to any value different than  $X$ . Thus operation  $\text{min\_in\_subtree1}_u(v)$  can be defined in the following way.

$\text{min\_in\_subtree1}_u(v) : \text{return } \text{min\_in\_subtree1}_r(\text{leader}(v))$ .

Operations  $\text{min\_in\_subtree2}_u$  and  $\text{find\_random\_X}_u$  have analogous definitions.

The implementations of operations  $\text{link}$  and  $\text{cut}$  on unrestricted dynamic trees are somewhat more involved. We give here only a detailed description of the im-

plementation of operation `cut`. A sequence of dynamic-tree operations on restricted trees which implements operation `cut(v)` on unrestricted trees can be derived from Figure 6. Node  $x$  is `leader(v)` and node  $y$  is `parentr(x)`. Nodes  $b$ ,  $c$ ,  $d$ , and  $e$  can be found using operations `parentr` and `childr`.

```

cutu(v):
  x ← leader(v);
  if x = v then v is the root of its (unrestricted) tree
  else
    y ← parentr(x); cutr(x); evertr(v);
    c ← parentr(x); b ← childr(x); d ← parentr(y); e ← childr(y);
    cutr(x); if b exists then begin cutr(b); linkr(b, c) end;
    cutr(y); if e exists then begin cutr(e); linkr(e, d) end;
    discard auxiliary nodes x and y.

```

The amortized time of one dynamic-tree operation on unrestricted trees is  $O(\log n)$ , if we implement operations `min_in_subtree1` or `find_random_X`, and  $O(q \log n)$ , if we implement operation `min_in_subtree2`.

#### 4. FLIP TREE

In this section we describe an application of dynamic trees with operation `find_random_X`. A degree- $d$  *flip tree* is a data structure which maintains an  $n$ -node undirected degree- $d$  tree  $T$  under a sequence of operations including operations of the following type.

`random_flip`:

Remove a random edge from the current tree  $T$ , splitting  $T$  into two trees  $T_1$  and  $T_2$ . Select two random nodes  $v$  and  $u$  of degree at most  $d - 1$ , one from tree  $T_1$  and the other from tree  $T_2$ . Using edge  $(v, u)$ , join trees  $T_1$  and  $T_2$  into a new degree- $d$  tree  $T'$ .

Such data structures arise in implementations of local search methods for degree-constrained minimum spanning tree problems. The notion of flip trees may be viewed as a generalisation of the *tour datatype*, defined and studied by Fredman, Johnson, McGeoch, and Ostheimer [1995]. The tour datatype maintains a traveling salesman tour and provides necessary operations for implementing local search methods for the traveling salesman problem. An “open” traveling salesman tour is a degree-2 spanning tree.

We implement a flip tree  $T$  by a collection of dynamic trees with operation `find_random_X`. We root tree  $T$  at an arbitrary node. The root keeps changing during the computation, but its exact position is never important outside of the data structure. Each node  $w$  of tree  $T$  has a key `X` or `not_X`, depending whether the degree of node  $w$  in tree  $T$  is less than  $d$  or equal to  $d$ . We need an additional dynamic-tree operation `degree(w)`, which returns the (undirected) degree of a node  $w$ . Operation `random_flip` can be implemented in the following way.

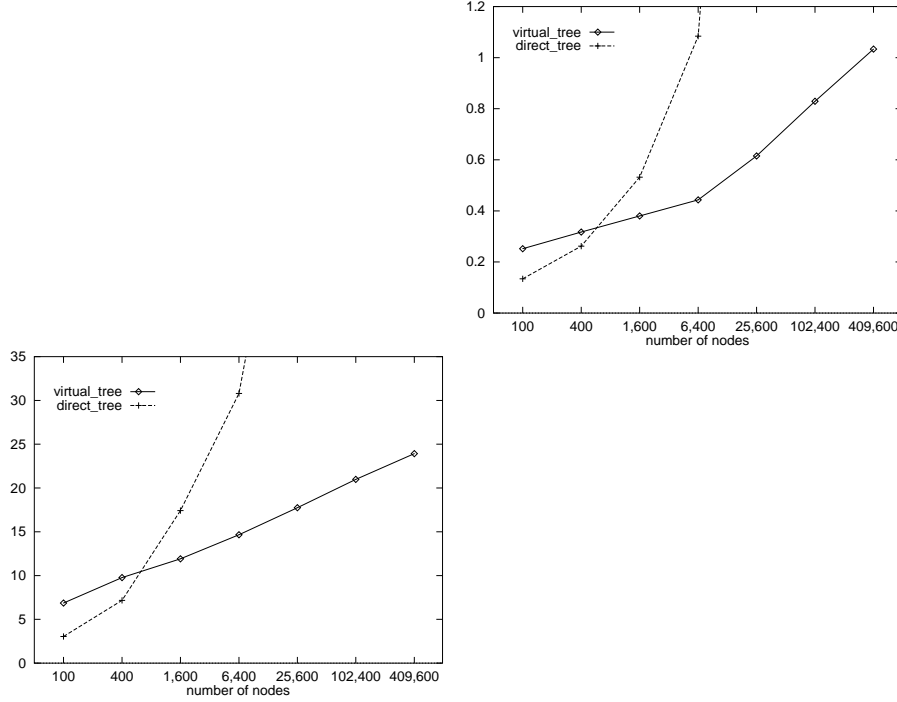


Fig. 7. A sequence of 10,000 operations `random_flip`, applied to an  $n$ -node tree obtained from “heap” by  $n$  operations `random_flip`. Left: Average (over 10 samples) execution times in seconds; for code `direct_tree` and the number of nodes 25,600 and 102,400, the times are 3.9 and 10.4, respectively. Right: Average depth of a node (for `direct_tree` the numbers are divided by 4).

```

random_flip:   $a \leftarrow$  a random node in  $T$  other than the root;
               $b \leftarrow$  parent( $a$ ); cut( $a$ ); new_key( $a$ , not_X); new_key( $b$ , not_X);
               $v \leftarrow$  find_random_X( $a$ ); evert( $b$ );  $u \leftarrow$  find_random_X( $b$ );
              evert( $v$ ); link( $v$ ,  $u$ );
              if degree( $v$ ) =  $d$  then new_key( $v$ , X);
              if degree( $u$ ) =  $d$  then new_key( $u$ , X).

```

We have coded two degree-3 flip-tree data structures. Code `virtual_tree` is based on the implementation of restricted dynamic trees with operation `find_random_X` described in Sections 2 and 3. Code `direct_tree` is based on a straightforward implementation of restricted dynamic trees: A dynamic tree  $D$  is represented internally by a tree which has the same structure as  $D$  (that is, no splaying). In code `virtual_tree` the amortized running time of one operation `random_flip` is  $O(\log n)$ . In code `direct_tree` the average running time of one operation `random_flip` is linear in the average depth of a node in  $D$ .

We compare the performance of codes `virtual_tree` and `direct_tree` in the following experiment. We generate an  $n$ -node degree-3 tree and then perform a sequence of operations `random_flip`. The initial tree is obtained by applying a sequence of  $n$  operations `random_flip` to either the heap-like tree (node  $i$  is connected by edges with nodes  $\lceil i/2 \rceil$ ,  $2i$ , and  $2i+1$ ) or the path-like tree (node  $i$  is connected by edges

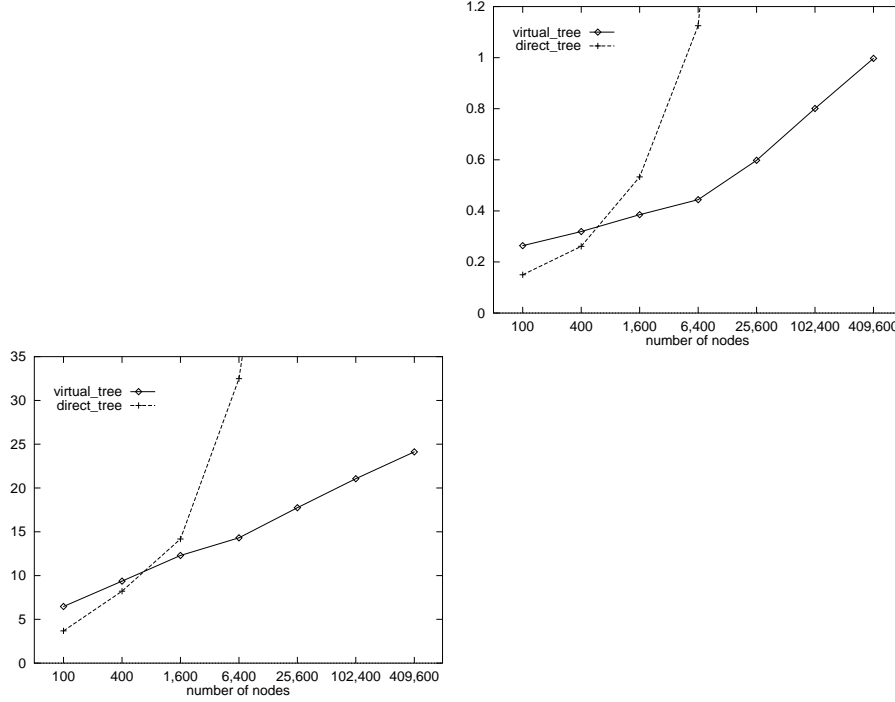


Fig. 8. A sequence of 10,000 operations `random_flip`, applied to an  $n$ -node tree obtained from “path” by  $n$  operations `random_flip`. Left: Average (over 10 samples) execution times in seconds; for code `direct_tree` and the number of nodes 25,600 and 102,400, the times are 3.4 and 8.7, respectively. Right: Average depth of a node (for `direct_tree` the numbers are divided by 4).

with nodes  $i - 1$  and  $i + 1$ ). In Figures 7 and 8 we plot average running times of sequences of 10,000 operations `random_flip`. We exponentially increase the number of nodes from  $10^2$  to  $4 \cdot 10^5$ . We also plot the average depth of a node in the tree representing the flip tree. These plots are consistent with the expectation that the average running time of one operation `random_flip` should be linearly related to the average depth of a node.

## 5. MINIMUM SPANNING FOREST IN DYNAMIC DENSE GRAPHS

Frederickson’s topology-trees data structure [Frederickson 1985] together with the sparsification method [Eppstein et al. 1992] give a data structure for maintaining the minimum spanning forest in a dynamic graph in  $O(\sqrt{n})$  worst-case time per one update – edge insertion or deletion. We denote by  $n$  and  $m$  the number of nodes and the number of edges in the underlying dynamic graph  $G = (V, E)$ . We assume that the set of nodes  $V$  is static, that is, it does not change during the computation. Recently Henzinger and King have improved this bound to  $O(n^{1/3} \log n)$  amortized time per one update [Henzinger and King 1997]. Amato, Cattaneo, and Italiano [1997] conducted an extensive empirical study of performance of several methods for maintaining the minimum spanning forest, including various variants of Frederickson’s data structure and the sparsification method, as well as the *ad hoc* method

code	edge insertions		edge deletions		random update (average)
	non-tree edge	tree edge	non-tree edge	tree edge	
mst1	Sleator-Tarjan dynamic trees		dynamic trees with operation <code>min_in_subtree1</code>		$\Theta\left(\frac{n^2}{m} \log n\right)$
	$O(\log n)$	$O(n \log n)$	$O(\log n)$	$O(n \log n)$	
mst2	Sleator-Tarjan dynamic trees		dynamic trees with operation <code>min_in_subtree2</code>		$\Theta\left(\frac{n^2}{m} \log n\right)$
	$O(\log n)$	$O(n \log n)$	$O(\log n)$	$O(n \log n)$	
mst_recompute	Sleator-Tarjan dynamic trees		recompute MST, if a tree edge deleted		$\Theta(n \log n)$
	$O(\log n)$	$O(\log n)$	$O(1)$	$O(m \log n)$	

Fig. 9. Asymptotic bounds for our dynamic minimum spanning forest codes.

of recomputing the whole minimum spanning forest whenever a tree edge, that is, an edge belonging to the minimum spanning forest, is deleted. Their experimental results show that their code `adhoc`, despite having  $O(m \log n)$  worst-case and  $O(m + n \log^2 n)$  average-case running time per one tree-edge deletion, outperforms their other codes, if the sequence of updates is random and the graph is not very sparse.

We have developed two codes, `mst1` and `mst2`, for the dynamic minimum spanning forest problem which are based on our implementations of dynamic trees with operations `min_in_subtree1` and `min_in_subtree2`, respectively. We have also developed a straightforward code `mst_recompute` which implements the *ad hoc* method. In all three codes insertions of edges are supported by the Sleator-Tarjan dynamic trees with operation `min_on_path` [Sleator and Tarjan 1985]. We denote this data structure by *ST*. In code `mst_recompute`, deletion of a tree edge is implemented by recomputing the whole minimum spanning forest using LEDA function `MIN_SPANNING_TREE`. The summary description of these three codes is shown in Figure 9. All three codes require  $\Theta(n^2)$  memory and are intended for dense graphs. Further details of codes `mst1` and `mst2` are presented in Sections 5.1 and 5.2, respectively. In Section 5.3 we discuss an experiment which we have conducted to compare the actual performance of all three codes.

### 5.1 Code `mst1`: Deletions of edges supported by $n$ dynamic tree collections

In code `mst1`, deletions of edges are supported by dynamic trees with operation `min_in_subtree1` in the following way. Let  $F$  denote the current minimum spanning forest. For each node  $x \in V$ , we maintain a data structure  $F(x)$ , which is a dynamic collection of trees with operation `min_in_subtree1` representing forest  $F$ . We maintain the invariant that the roots in all collections  $F(x)$  are the same; that is, a given node is either a root in all collections  $F(x)$  or in none of them. For each pair of nodes  $v$  and  $w$ , the key of node  $v$  in collection  $F(w)$  and the key of node  $w$  in collection  $F(v)$  are the same and are equal to the weight of edge  $(v, w)$ , or to

```

delete_edge(v, u):
  C ← ∅; — the set of candidate edges for replacing edge (v, u) —
  new_key(F(v), u, ∞); new_key(F(u), v, ∞);
  if (v, u) is a tree edge then
    delete edge (v, u) from ST;
    — note: the directions of edges are the same in all collections F(x), x ∈ V —
    if parent(F(v), u) = v then begin temp ← v; v ← u; u ← temp end;
    r ← find_root(F(v), u);
    for each x ∈ V do
      cut(F(x), v);
      — note: node find_root(F(x), u) is node r —
      if find_root(F(x), x) = r then C ← C ∪ {(x, min_in_subtree1(F(x), v))}
    if C ≠ ∅ then
      (p, q) ← a minimum weight edge in C; add edge (p, q) to ST;
      for each x ∈ V do begin evert(F(x), p); link(F(x), p, q); end.

```

Fig. 10. Implementation of the edge-deletion operation in code `mst1`

infinity, if  $(v, w)$  is not an edge in the current graph.

The implementation of the edge-deletion operation is shown in Figure 10 (the details of the updates of data structure  $ST$  are omitted). For each dynamic-tree operation, we introduce an additional parameter indicating the dynamic-tree collection which is to be used. To delete an edge  $(v, u)$  from the graph, we first set to infinity the appropriate keys in collections  $F(v)$  and  $F(u)$  and check whether edge  $(v, u)$  is a tree edge. If  $(v, u)$  is a tree edge, then we perform the following updates. We delete edge  $(v, u)$  from data structure  $ST$  and find a replacement edge using collections  $F(x)$ ,  $x \in V$ . Let the direction of edge  $(v, u)$  in collections  $F(x)$ ,  $x \in V$ , be from  $v$  (child) to  $u$  (parent). The roots in all these collections are the same, so the directions of edges are also the same. For each node  $x \in V$ , we delete edge  $(v, u)$  from collection  $F(x)$  by executing operation  $\text{cut}(F(x), v)$ . The removal of edge  $(v, u)$  from forest  $F$  splits one tree of  $F$  into two trees  $T_v$  and  $T_u$ . Tree  $T_v$  contains node  $v$  and tree  $T_u$  contains node  $u$ . We define a set of edges  $C$  in the following way,

$$C = \{(x, \text{min\_in\_subtree1}(F(x), v)) : \text{find\_root}(F(x), x) = \text{find\_root}(F(x), u)\}.$$

Thus set  $C$  contains a minimum-weight edge joining node  $x$  and tree  $T_v$ , for each node  $x$  in tree  $T_u$ . A minimum-weight edge  $(p, q)$  from  $C$  replaces edge  $(v, u)$  in forest  $F$ . Using operations  $\text{evert}$  and  $\text{link}$ , we add edge  $(p, q)$  to  $F(x)$ , for each  $x \in V$ , and to  $ST$ .

Each insertion and deletion of a non-tree edge is implemented by a constant number of dynamic tree operations, so it runs in  $O(\log n)$  amortized time. Each insertion and deletion of a tree edge involves  $\Theta(n)$  dynamic-tree operations (a constant number of dynamic-tree operations in  $ST$  and in each collection  $F(x)$ ,  $x \in V$ ) so it runs in  $O(n \log n)$  amortized time.

## 5.2 Code `mst2`: Deletions of edges supported by one dynamic tree collection

In code `mst1`, deletions of edges are supported by  $n$  separate dynamic tree collections, so a substantial part of the run time is always spent on restructuring virtual



```

delete_edge(v, u):
  new_key(DT, v, u, ∞); new_key(DT, u, v, ∞);
  if (v, u) is a tree edge then
    delete edge (v, u) from ST;
    evert(DT, u); cut(DT, v);
    P – an array indexed by nodes;
    for each x ∈ V do
      if find_root(ST, x) = find_root(ST, u) then P[x] = 1 else P[x] = 0;
    (p, q) ← min_in_subtree2(DT, v, P);
    if key(DT, p, q) < ∞ then
      evert(DT, p); link(DT, p, q); add (p, q) to ST.

```

Fig. 11. Implementation of the edge-deletion operation in code `mst2`

trees. To cut down on this restructuring, we develop code `mst2`, in which deletions of edges are based on representing the current minimum spanning forest  $F$  by one dynamic tree collection with operation `min_in_subtree2`. Let  $DT$  denote this collection. Each node in  $DT$  has an array of  $n$  keys indexed by nodes. The value  $\text{key}(DT, x, y)$  of the key at node  $x$  corresponding to node  $y$  is equal to the weight of edge  $(x, y)$ , or to infinity, if  $(x, y)$  is not an edge in the current graph.

The implementation of the deletion of a graph edge  $(v, u)$  is shown in Figure 11. We first set keys  $\text{key}(DT, v, u)$  and  $\text{key}(DT, u, v)$  to infinity, and check whether edge  $(v, u)$  is a tree edge. If  $(v, u)$  is a tree edge, then we perform the following updates. We delete edge  $(v, u)$  from  $ST$  and  $DT$ . Deleting edge  $(v, u)$  from  $DT$  is done by executing operations `evert(DT, u)` and `cut(DT, v)`. Let  $T_v$  and  $T_u$  be the new trees in forest  $F$  obtained by the removal of edge  $(v, u)$ ; tree  $T_v$  is the one which contains node  $v$ . Let  $P$  be an array of size  $n$  indexed by the nodes in  $V$  such that for each node  $x \in V$ ,  $P[x] = 1$  if and only if  $x$  belongs to tree  $T_u$ . The tree in collection  $DT$  representing tree  $T_v$  is rooted at node  $v$ . Therefore, operation `min_in_subtree2(DT, v, P)` returns a pair  $(p, q)$  of nodes which minimises  $\text{key}(DT, p', q')$  over all nodes  $p'$  which belong to tree  $T_v$  and all nodes  $q'$  which belong to tree  $T_u$ . If the returned pair is an edge in the current graph, then this edge replaces edge  $(v, u)$  in the minimum spanning forest  $F$ .

One dynamic-tree operation on collection  $DT$  takes  $O(n \log n)$  amortized time. Each insertion and deletion of a tree edge involves a constant number of dynamic-tree operations on collection  $DT$ , so it takes  $O(n \log n)$  amortized time. Each insertion and deletion of a non-tree edge involves two operations `new_key` on collection  $DT$ . In the virtual-tree representation of dynamic trees, changing the value of a key at a node  $v$  is normally done by applying operation `splay(v)` and then performing necessary updates of the key-related attributes of node  $v$ . If we followed this approach, the amortize time of one operation `new_key` on collection  $DT$ , and consequently the amortized time of one update involving a non-tree edge, would be  $O(n \log n)$ . To achieve a  $O(\log n)$  bound, we apply operation `splay(v)` in operation `new_key` only if the path in the virtual tree from node  $v$  to the root is short.

We implement operation `new_key(DT, x, y, k)` (set the key of node  $x$  with index  $y$  to value  $k$ ) in the following way. We change the value of  $\text{key}(DT, x, y)$  to value  $k$ , and then follow the path in the virtual tree from node  $x$  to the root, updating values

$\text{min\_key}(DT, z, y)$  at each encountered node  $z$ . If this path contains more than  $\gamma \log n$  nodes, where  $\gamma$  is some fixed constant, then we apply operation  $\text{splay}(x)$ . The running time of a cheap operation  $\text{new\_key}$  (an operation without splaying) is  $O(\log n)$ , while the running time of an expensive operation  $\text{new\_key}$  is  $\Theta(\delta n)$ , where  $\delta$  is the length of the path traversed in the virtual tree;  $\delta \geq \gamma \log n$ . Next we show that the amortized time of one operation  $\text{new\_key}$  is  $O(\log n)$ .

Consider a sequence  $\sigma$  of dynamic-tree operations on collection  $DT$ , assuming that the collection is initially empty. Let  $q_1$  denote the number of expensive operations  $\text{new\_key}$  in sequence  $\sigma$ , and let  $q_2$  denote the number of operations in sequence  $\sigma$  other than operations  $\text{new\_key}$ . The running time of each cheap operation  $\text{new\_key}$  is  $O(\log n)$ , and the running time of each operation in sequence  $\sigma$  other than a cheap operation  $\text{new\_key}$ , amortized over all such operations, is  $O(n \log n)$  (see Section 3.2), that is, at most  $\beta n \log n$ , for some constant  $\beta$ . Select constant  $\gamma$  such that the actual time of one expensive operation  $\text{new\_key}$  is at least  $2\beta n \log n$ . Thus the total time spent on the expensive operations  $\text{new\_key}$  in sequence  $\sigma$  is at least  $2q_1 \beta n \log n$ , but cannot be greater than  $(q_1 + q_2) \beta n \log n$ . Therefore  $q_1 \leq q_2$ , and we can distribute the cost of the expensive operations  $\text{new\_key}$  over the operations other than operations  $\text{new\_key}$ . This way a cheap operation  $\text{new\_key}$ , an expensive operation  $\text{new\_key}$ , and any other operation in sequence  $\sigma$  have amortized running times  $O(\log n)$ , 0, and  $O(n \log n)$ , respectively. Hence the amortized running time of one operation  $\text{new\_key}$  on collection  $DT$  is  $O(\log n)$ , so the amortized running time of one update of a non-tree edge in code `mst2` is  $O(\log n)$ .

### 5.3 Actual performance of our dynamic minimum spanning tree codes

We use our dynamic minimum spanning tree codes `mst1`, `mst2`, and `mst_recompute` in the following experiment. We generate a random undirected weighted graph  $G = (V, E)$  with  $n$  nodes and  $n(n-1)/4$  edges, and then maintain the spanning tree in  $G$  under a random sequence of deletions and insertions of edges. Each type of update (deletion or insertion) has the same probability and edges for deletion and insertion are selected uniformly at random. The weights of the inserted edges are also chosen at random.

The average running times of all three codes are shown in Figure 12. The theoretical analysis predicts that the average running time of one update in such experiment should be  $\Theta(\log n)$  for codes `mst1` and `mst2`, and  $\Theta(n)$  for code `mst_recompute`. Our experimental results seem to be consistent with this prediction. Code `mst2` is about 2.5-3 times faster than code `mst1`. This difference is consistent with the ratio of the average number of updates of the node pointers to the average number of updates of the node attributes  $\text{min\_key}$  in one operation  $\text{splay}$  in code `mst1`, which is roughly equal to 1.5. In code `mst2`, the number of updates of the node pointers becomes insignificant, while the number of updates of attributes  $\text{min\_key}$  remains more-less the same as in code `mst1`. Code `mst_recompute` is based on the same method as code `adhoc` developed by Amato, Cattaneo, and Italiano [1997]. Since code `adhoc` performed quite well in their experiments and our code `mst2` performs clearly better than code `mst_recompute` in our experiments, code `mst2` may actually be quite competitive in practice.




Fig. 12. Maintaining the minimum spanning tree in dynamic dense graphs. The initial density is  $1/2$ . Edge weights are random numbers from  $[0, 10000]$ . The times are averages over 10 samples.

## 6. CONCLUSIONS AND FURTHER WORK

We have developed a generic implementation of dynamic trees which are to support “in-subtree” operations. Our implementation is based on Sleator and Tarjan’s dynamic-trees data structure which uses splay trees [Sleator and Tarjan 1985]. Using our generic implementation, we implement three variants of dynamic trees, which support “in-subtree” operations `find_random_X`, `min_in_subtree1`, and `min_in_subtree2`, respectively. We use the first variant to implement a flip-tree data structure and the other two to implement two dynamic minimum spanning tree data structures. Our experimental results suggest that in some applications our dynamic-tree implementations may lead to substantially faster codes than codes based on more straightforward approaches.

One can think about a number of potential experiments which may provide further insight into practical value of our dynamic-trees implementations. For example, it would be interesting to evaluate the performance of our dynamic trees with operation `find_random_X` in the simulated-annealing method for degree-constrained minimum spanning tree problems. The next step in the evaluation of practical usefulness of our dynamic trees with operations `min_in_subtree1` and `min_in_subtree2` in the context of the dynamic minimum spanning tree problem could be a direct comparison of the performance of our codes `mst1` and `mst2` with the performance of codes which are based on theoretically fast methods.

Codes `mst1` and `mst2` are only practical for dense graphs. Could dynamic trees with operations `min_in_subtree1` and `min_in_subtree2` lead to data structures for the dynamic minimum spanning tree problem which are practical also for sparse graphs? It seems that to maintain efficiently the minimum spanning tree in a sparse graph, one needs to maintain a hierarchical decomposition of the tree. Observe that the

virtual-tree representation of a dynamic tree  $T$  can be actually viewed as a hierarchical decomposition of  $T$ . We believe that this decomposition can be further exploited.

Our flip-tree data structure (code `virtual_tree` in Section 4) can be viewed as a generalisation of the dynamic traveling-salesman data structure based on splay trees. Fredman, Johnson, McGeoch, and Ostheimer [1995] describe this data structure and show that it gives good performance of local-search optimisation algorithms for TSP problem. They also show, however, that some other dynamic traveling-salesman data structures, for example their data structure based on two-level trees, are usually faster. It would be interesting to design and evaluate a generic dynamic-tree data structure which generalises the two-level-tree data structure.

## REFERENCES

- AHUJA, R. K., ORLIN, J. B., AND TARJAN, R. E. 1989. Improved Time Bounds for the Maximum Flow Problem. *SIAM J. Comput.* 18, 939–954.
- AMATO, G., CATTANEO, G., AND ITALIANO, G. F. 1997. Experimental analysis of dynamic minimum spanning tree algorithms. In *Proc. 8th Annual ACM-SIAM Symposium on Discrete Algorithms* (1997).
- BENT, S. W., SLEATOR, D. D., AND TARJAN, R. E. 1985. Biased Search Trees. *SIAM J. Comput.* 14, 3, 545–568.
- CERNY, V. 1985. A thermodynamical approach to the traveling salesman problem: An efficient simulation algorithm. *J. Optimization Theory and Appl.* 45, 41–51.
- EPPSTEIN, D., GALIL, Z., ITALIANO, G. F., AND NISSENZWEIG, A. 1992. Sparsification – a technique for speeding up dynamic graph algorithms. In *Proc. 33rd IEEE Annual Symposium on Foundations of Computer Science* (1992), pp. 60–69.
- FREDERICKSON, G. N. 1985. Data structures for on-line updateing of minimum spanning trees, with applications. *SIAM J. Comput.* 14, 4, 781–798.
- FREDMAN, M. L., JOHNSON, D. S., MCGEOCH, L. A., AND OSTHEIMER, G. 1995. Data structures for traveling salesmen. *J. Alg.* 18, 3, 432–479.
- GOLDBERG, A. V., TARDOS, E., AND TARJAN, R. E. 1990. Network Flow Algorithms. In B. KORTE, L. LOVÁS, H. PRÖMEL, AND A. SCHRIJVER Eds., *Paths, Flows, and VLSI-layout*, pp. 101–164. Springer-Verlag.
- GOLDBERG, A. V. AND TARJAN, R. E. 1989. Finding minimum-cost circulations by canceling negative cycles. *J. Assoc. Comput. Mach.* 36, 388–397.
- GOLDBERG, A. V. AND TARJAN, R. E. 1990. Finding Minimum-Cost Circulations by Successive Approximation. *Math. Oper. Res.* 15, 430–466.
- HENZINGER, M. R. AND KING, V. 1997. Maintaining minimum spanning trees in dynamic graphs. In *Proc. 24th International Colloquium on Automata, Languages and Programming* (July 1997).
- KIRKPATRICK, S., C. D. GELATT, J., AND VECCHI, M. P. 1983. Optimization by simulated annealing. *Science* 220, 671–680.
- MEHLHORN, K. AND NÁHER, S. 1995. LEDA, A platform for combinatorial and geometric computing. *Comm. ACM* 38, 1, 96–102.
- SLEATOR, D. D. AND TARJAN, R. E. 1983. A Data Structure for Dynamic Trees. *J. Comput. System Sci.* 26, 362–391.
- SLEATOR, D. D. AND TARJAN, R. E. 1985. Self-adjusting binary search trees. *J. Assoc. Comput. Mach.* 32, 652–686.
- TARJAN, R. E. 1983. *Data Structures and Network Algorithms*. Society for Industrial and Applied Mathematics, Philadelphia, PA.